

Task Report: Twitter Text Classification

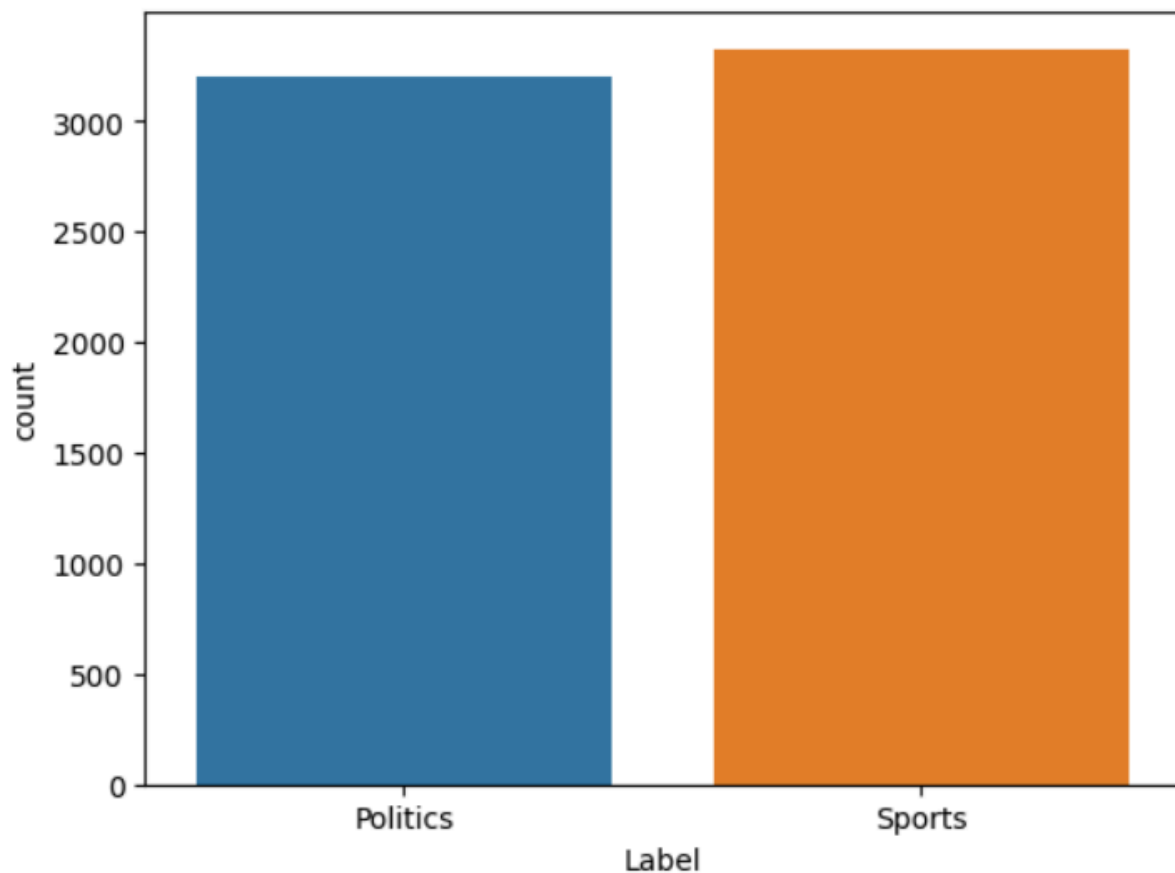
I. Data Exploration:

1. Data Loading and Visualization:

- Loaded the training and test datasets using `pd.read_csv`:

```
train_data = pd.read_csv('/kaggle/input/deeptweets/train.csv')  
test_data = pd.read_csv('/kaggle/input/deeptweets/test.csv')
```

- Utilized Seaborn to visualize the distribution of labels in the training set using `sns.countplot`:





WordCloud for "Sports"



2. Text Processing:

- Defined a text processing function ``Text_Processing`` to clean and preprocess the tweet text.
- Converted text to lowercase, removed special characters, and URLs using regular expressions.
- Tokenized the text using NLTK's ``word_tokenize``.
- Applied stemming using the PorterStemmer to reduce words to their root form.
- Created a new column ``ProcessedText`` in the training dataset to store the preprocessed text.

```

# Initialization of the stemmer
stemmer = PorterStemmer()

def Text_Processing(text):
    text = text.lower()
    # Remove special characters and URLs
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'^a-zA-Z\s', '', text)
    # Tokenize the text
    text = word_tokenize(text)
    # Stemming words
    text = [stemmer.stem(word) for word in text]
    # Join the processed words back into a sentence
    text = ' '.join(text)
    return text

```

III. Data Splitting:

- Split the dataset into training and testing sets using `train_test_split`:

```

# Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

IV. Model Building:

1. Logistic Regression Pipeline:

- Constructed a machine learning pipeline using `Pipeline` from scikit-learn.
- Utilized a bag-of-words representation with `CountVectorizer` (unigrams and bigrams).
- Employed a logistic regression classifier (`LogisticRegression`) as the predictive model.

```

LR = Pipeline([
    ('bag_of_word' , CountVectorizer(ngram_range=(1,2))),
    ('LR' , LogisticRegression())
])
LR.fit(X_train, y_train)
y_pred = LR.predict(X_test)
print(classification_report(y_test, y_pred))

```

2. Model Training and Evaluation:

- Fitted the pipeline on the training data (`X_train`, `y_train`).
- Made predictions on the test data and printed a classification report using `classification_report`.
- Displayed a confusion matrix to visualize model performance.

V. Results and Analysis:

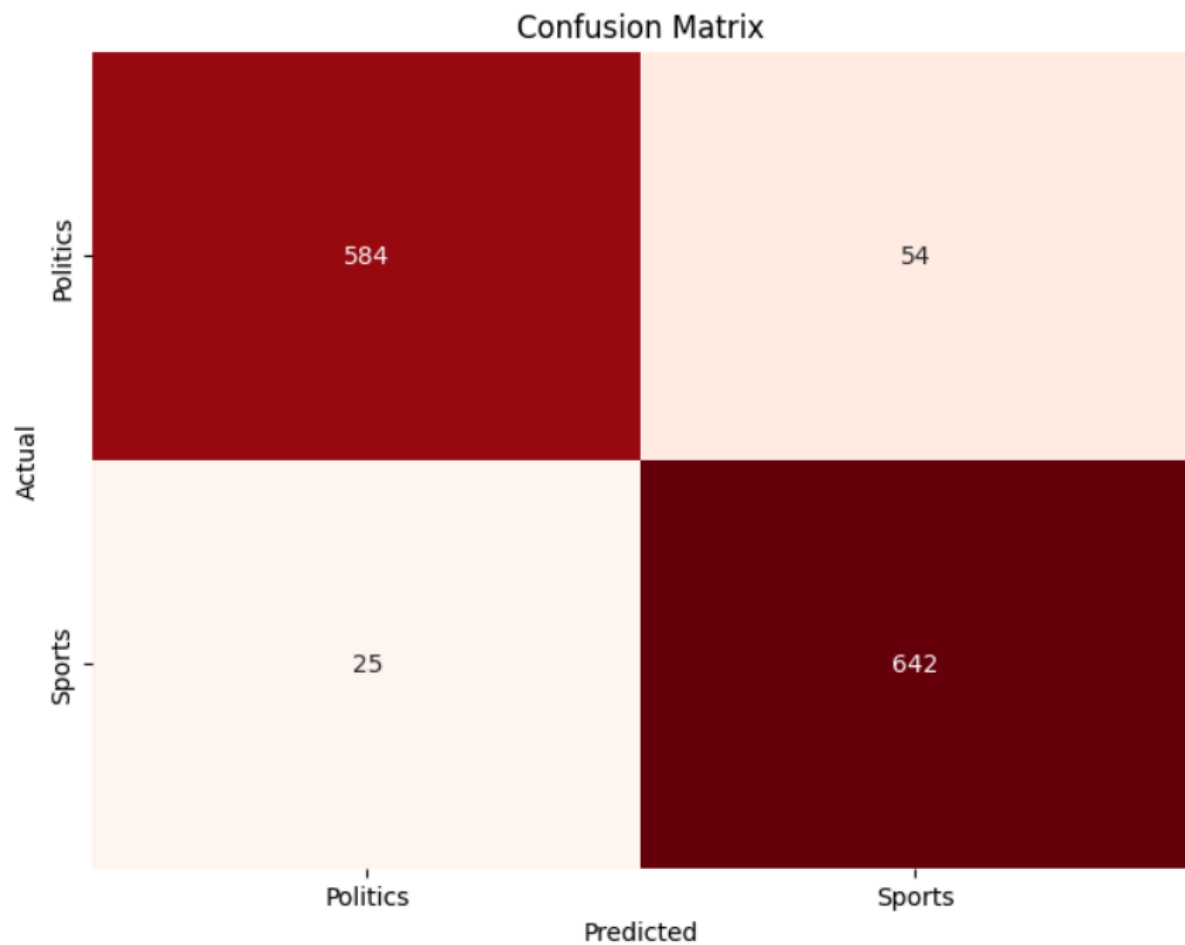
1. Classification Report:

- Examined precision, recall, and F1-score for both "Politics" and "Sports" labels.
- Assessed the overall model performance on the test set.

	precision	recall	f1-score	support
Politics	0.96	0.92	0.94	638
Sports	0.92	0.96	0.94	667
accuracy			0.94	1305
macro avg	0.94	0.94	0.94	1305
weighted avg	0.94	0.94	0.94	1305

2. Confusion Matrix Visualization:

- Visualized the confusion matrix using Seaborn's heatmap to understand model predictions.
- The heatmap displays actual vs. predicted labels.



VI. Conclusion:

- Summarized key findings, including model performance metrics.
- Provided insights into the effectiveness of the chosen model and preprocessing techniques.
- Suggested potential areas for improvement, such as experimenting with different vectorization techniques or exploring other classification algorithms.